

HAPPYMONK TASK

Barath Vignesh.S
barathvignesh01@gmail.com

TASK 1

QUESTION 1

1. Consider a large dataset (say, a time series) A. Also, consider a smaller dataset B. How do you ensure whether sets A and B identify the same variable? Illustrate it with a Python script.

DATASET A - LARGE DATASET

The aim is to recommend movies on amazon prime based on voting and ratings given to each movie. The task is completed by taking 2 large dataset and combining them accordingly .

It consists of field like movie_id , title , cast ,crew ,budget ,keywords ,original_language, runtime, original_title , overview, popularity, production_companies, production_countries ,release_date, revenue, spoken_languages, status, tagline, vote_average, vote_count.

	A	B	C			
1	movie_id	title	cast	crew		
2	19995	Avatar	[{"cast_id": 242, "credit_id": "52fe48009c"}, {"cast_id": 4, "credit_id": "52fe4232c3"}, {"cast_id": 285, "credit_id": "52fe479ac3"}, {"cast_id": 1, "credit_id": "54805967c3"}, {"cast_id": 2, "credit_id": "52fe4781c3"}, {"cast_id": 5, "credit_id": "52fe479ac3"}, {"cast_id": 30, "credit_id": "52fe4252c3"}, {"cast_id": 34, "credit_id": "52fe46db9c3"}, {"cast_id": 1, "credit_id": "55d5f7d4c3"}, {"cast_id": 3, "credit_id": "52fe4273c3"}, {"cast_id": 76, "credit_id": "52fe4273c3"}, {"cast_id": 18, "credit_id": "553bf2369c3"}, {"cast_id": 3, "credit_id": "553bef6a9c3"}, {"cast_id": 1, "credit_id": "52fe43b29c3"}, {"cast_id": 37, "credit_id": "52fe4211c3"}, {"cast_id": 4, "credit_id": "52fe4928c3"}, {"cast_id": 2, "credit_id": "52fe4799c3"}, {"cast_id": 1, "credit_id": "55a239e69c3"}, {"cast_id": 3, "credit_id": "52fe4495c3"}, {"cast_id": 15, "credit_id": "566b4f54c3"}, {"cast_id": 4, "credit_id": "52fe45b7c3"}, {"cast_id": 10, "credit_id": "548ad49a9c3"}, {"cast_id": 56, "credit_id": "5395a60dc3"}, {"cast_id": 1, "credit_id": "52fe43f2c3"}, {"cast_id": 3, "credit_id": "52fe4926c3"}, {"cast_id": 43, "credit_id": "52fe4348c3"}, {"cast_id": 5, "credit_id": "52fe422ec3"}, {"cast_id": 254, "credit_id": "52fe422ec3"}]			
3						
4						
5						
6						
7						
8						
9						
10						
11						
12						
13						
14						
15						
16						
17						
18						
19						
20						
21						
22						
23						
24						
25						
26						

By taking the vote_count and vote_average into account , calculated the weighted_rating by using the IMDB formula

```
[ ] def weighted_rating(x, m=m, C=C):
    v = x['vote_count']
    R = x['vote_average']
    return (v/(v+m) * R) + (m/(m+v) * C)
```

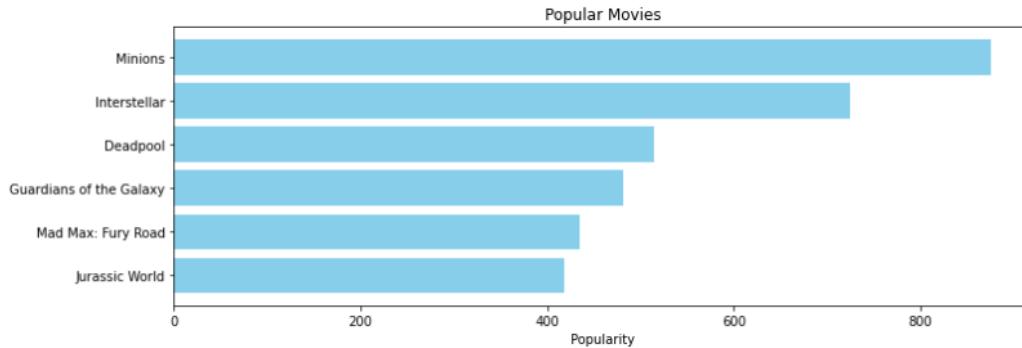
Created a new feature score and calculated the weighted rating value . Sorted the movies on scores and printed the top movies

```
[ ] q_movies = q_movies.sort_values('score', ascending=False)
q_movies[['title', 'vote_count', 'vote_average', 'score']].head(10)
```

	title	vote_count	vote_average	score
1881	The Shawshank Redemption	8205	8.5	8.059258
662	Fight Club	9413	8.3	7.939256
65	The Dark Knight	12002	8.2	7.920020
3232	Pulp Fiction	8428	8.3	7.904645
96	Inception	13752	8.1	7.863239
3337	The Godfather	5893	8.4	7.851236
95	Interstellar	10867	8.1	7.809479
809	Forrest Gump	7927	8.2	7.803188
329	The Lord of the Rings: The Return of the King	8064	8.1	7.727243
1990	The Empire Strikes Back	5879	8.2	7.697884

BAR CHART:

```
Text(0.5, 1.0, 'Popular Movies')
```



Preprocessing involves removal of stop words like 'the' , 'a'. Replaced NaN with empty string.

```
[ ] from sklearn.feature_extraction.text import TfidfVectorizer
tfidf = TfidfVectorizer(stop_words='english')

df2['overview'] = df2['overview'].fillna('')
tfidf_matrix = tfidf.fit_transform(df2['overview'])

tfidf_matrix.shape
```

(4803, 20978)

Created a function that takes in movie genres and matches with the scores .The cosine similarity matches the similar movies and displays the top movies based on the scores that are higher.

```
get_recommendations('The Dark Knight Rises')
```

```
65                      The Dark Knight
299                     Batman Forever
428                     Batman Returns
1359                    Batman
3854    Batman: The Dark Knight Returns, Part 2
119                      Batman Begins
2507                     Slow Burn
9           Batman v Superman: Dawn of Justice
1181                     JFK
210                      Batman & Robin
Name: title, dtype: object
```

DATASET 2 - SMALL DATASET

A	B	C
movielid	title	genres
1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy
2	Jumanji (1995)	Adventure Children Fantasy
3	Grumpier Old Men (1995)	Comedy Romance
4	Waiting to Exhale (1995)	Comedy Drama Romance
5	Father of the Bride Part II (1995)	Comedy
6	Heat (1995)	Action Crime Thriller
7	Sabrina (1995)	Comedy Romance
8	Tom and Huck (1995)	Adventure Children
9	Sudden Death (1995)	Action
10	GoldenEye (1995)	Action Adventure Thriller
11	American President, The (1995)	Comedy Drama Romance
12	Dracula: Dead and Loving It (1995)	Comedy Horror
13	Balto (1995)	Adventure Animation Children
14	Nixon (1995)	Drama
15	Cutthroat Island (1995)	Action Adventure Romance
16	Casino (1995)	Crime Drama
17	Sense and Sensibility (1995)	Drama Romance
18	Four Rooms (1995)	Comedy
19	Ace Ventura: When Nature Calls (1995)	Comedy
20	Money Train (1995)	Action Comedy Crime Drama Thriller
21	Get Shorty (1995)	Comedy Crime Thriller

Since some movies don't have the information about the year in the column , we need to handle them . I have changed the column name from title to title year , removed leading and ending whitespaces in the title_year and created separate columns for title and year.



```
def extract_title(title):
    year = title[len(title)-5:len(title)-1]

    if year.isnumeric():
        title_no_year = title[:len(title)-7]
        return title_no_year
    else:
        return title

def extract_year(title):
    year = title[len(title)-5:len(title)-1]
    if year.isnumeric():
        return int(year)
    else:
        return np.nan

movies.rename(columns={'title':'title_year'}, inplace=True)
movies['title_year'] = movies['title_year'].apply(lambda x: x.strip())
movies['title'] = movies['title_year'].apply(extract_title)
movies['year'] = movies['title_year'].apply(extract_year)
```

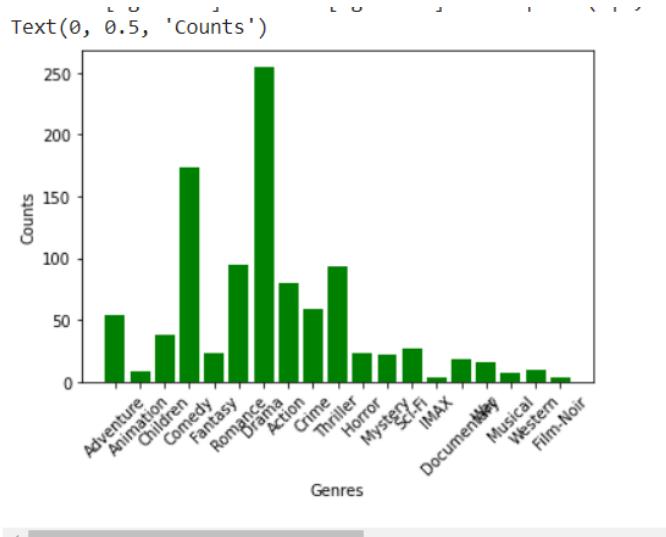
Removed the movies without genre info and resetted the index , Then removed the ‘|’ in the genre column and counted the number of occurrence for each genre in dataset

```

movies['genres'] = movies['genres'].str.replace('|', ' ')
counts = dict()
for i in movies.index:
    for g in movies.loc[i,'genres'].split(' '):
        if g not in counts:
            counts[g] = 1
        else:
            counts[g] = counts[g] + 1

```

BARCHART :



By using TfidfVectorizer, each feature index in the column is given with a unique token id . Builded the cosine similarity matrix and found the similarity between them.

```

print(list(enumerate(tfidf_vector.get_feature_names())))
[(*, 'action'), (*, 'adventure'), (*, 'animation'), (*, 'children'), (*, 'comedy'),
/usr/local/lib/python3.8/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning
warnings.warn(msg, category=FutureWarning)

```

```

print(sim_matrix)

[[1.          0.75572094 0.15025974 ... 0.09574404 0.          0.24605568]
 [0.75572094 1.          0.          ... 0.          0.          0.          ]
 [0.15025974 0.          1.          ... 0.23762251 0.67115372 0.61067374]
 ...
 [0.09574404 0.          0.23762251 ... 1.          0.          0.38911532]
 [0.          0.          0.67115372 ... 0.          1.          0.          ]
 [0.24605568 0.          0.61067374 ... 0.38911532 0.          1.          ]]

```

Fuzzywuzzy ,the similar matrices are grouped together . `find_closest_title` is used to return the most similar title to the word that is given as input.

```

def find_closest_title(title):
    leven_scores = list(enumerate(movies['title'].apply(matching_score, b=title)))
    sorted_leven_scores = sorted(leven_scores, key=lambda x: x[1], reverse=True)
    closest_title = get_title_from_index(sorted_leven_scores[0][0])
    distance_score = sorted_leven_scores[0][1]
    return closest_title, distance_score

```

Finally , the distance score is set to a threshold value of 100 , it even calculates correctly when the user makes a mistake in typing . And removed the typed movie itself . Then the list of movies is displayed and we can alter the no of input movies

```

contents_based_recommender('Monsters, Inc.', 20)

Did you mean Dunston Checks In?

Here's the list of movies similar to Dunston Checks In.

It Takes Two (1995)
Big Green, The (1995)
Heavyweights (Heavy Weights) (1995)
Richie Rich (1994)
Baby-Sitters Club, The (1995)
Now and Then (1995)
Babe (1995)
White Balloon, The (Badkonake sefid) (1995)
Fluke (1995)
Little Princess, A (1995)
Gordy (1995)
Flintstones, The (1994)
Addams Family Values (1993)
Mighty Morphin Power Rangers: The Movie (1995)
Tom and Huck (1995)
Amazing Panda Adventure, The (1995)
Casper (1995)
Far From Home: The Adventures of Yellow Dog (1995)
Lassie (1994)
Free Willy 2: The Adventure Home (1995)

```

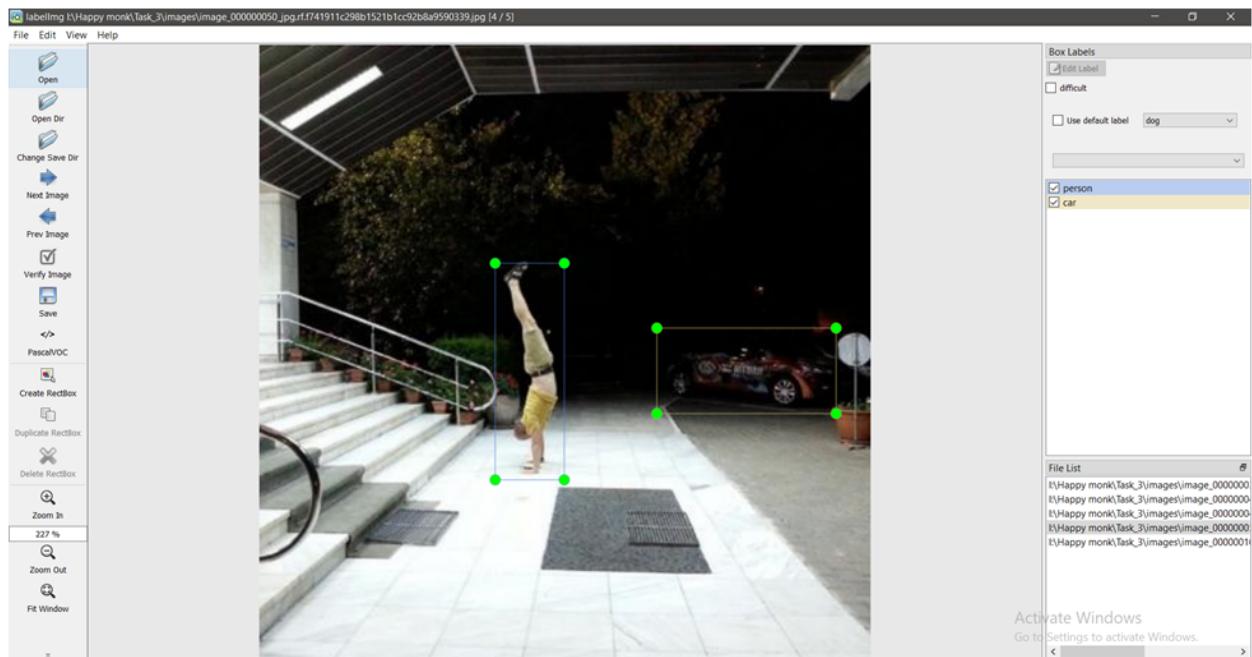
TASK 2

QUESTION 2

2. Collect data (images) and annotate them for two classes: Person and vehicle. You may use platforms such as LabelImg for annotations. You may limit to 800 images for the dataset.

Perform object detection on your collected dataset and find the mean distance between the two classes in each image. You may use YOLOv5 for detection.

Annotating using LabelImg:



Image_name.xml file for yolo model:

```
image_00000050.jpg.rf.f741911c298b1521b1cc92b8a9590339 - Notepad
File Edit Format View Help
1 0.441106 0.532452 0.112981 0.353365
4 0.795673 0.531250 0.293269 0.139423
```

Using Pre Trained YOLOV5 model - COCO model:

yolov5 prerequisites

Yolov5 Prerequisite

```
]:  
!git clone https://github.com/ultralytics/yolov5 # clone  
!cd yolov5  
!pip install -qr requirements.txt # install  
  
Cloning into 'yolov5'...  
remote: Enumerating objects: 14461, done.  
remote: Counting objects: 100% (90/90), done.  
remote: Compressing objects: 100% (64/64), done.  
remote: Total 14461 (delta 56), reused 53 (delta 25), pack-reused 14371  
Receiving objects: 100% (14461/14461), 13.57 MiB | 9.62 MiB/s, done.  
Resolving deltas: 100% (9957/9957), done.  
/kaggle/working/yolov5  
WARNING: Running pip as the 'root' user can result in broken permissions and conflicting behaviour with the system package manager. It is recommended  
to use a virtual environment instead: https://pip.pypa.io/warnings/venv  
..
```

Importing the necessary libraries and fitting the pretrained model:

Import Libraries

```
]:  
import torch  
from matplotlib import pyplot as plt  
import numpy as np  
import cv2
```

Loading Pretrained Model

```
]:  
model = torch.hub.load('ultralytics/yolov5', 'yolov5s')  
model.classes = [0, 2] # person and car  
  
  
Downloading: "https://github.com/ultralytics/yolov5/archive/master.zip" to /root/.cache/torch/hub/master.zip  
YOLOv5 🚀 2022-12-11 Python-3.7.12 torch-1.11.0+cpu CPU  
  
Downloading https://github.com/ultralytics/yolov5/releases/download/v7.0/yolov5s.pt to yolov5s.pt...  
0%|          | 0.00/14.1M [00:00, ?B/s]  
Fusing layers...  
YOLOv5s summary: 213 layers, 7225885 parameters, 0 gradients  
Adding AutoShape...
```

Object Detection on Sample Image (Car - Person):

Pretrained-Model on Sample Data

```
i]:  
    img = 'I:/car-person-v2-roboflow/Car-Person-v2-Roboflow-Owais-Ahmad/test/images/image_00000102.jpg.rf.79e39dcbcd8179ea9e837ce05fb096a1.jpg'  
  
    # Inference  
    results = model(img)  
    results.print() # or .show(), .save()  
  
  
image 1/1: 416x416 5 persons, 2 cars  
Speed: 72.7ms pre-process, 360.1ms inference, 15.9ms NMS per image at shape (1, 3, 640, 640)  
  
[]:  
    %matplotlib inline  
    plt.imshow(np.squeeze(results.render()))  
    plt.show()  
  
  
Activ
```

Calculating mean distance between classes:

The midpoints of the objects and mean difference of two points are calculated

Calculating Mean Difference

```
[]:  
    res = results.pandas().xyxy[0]  
    rn = {'car':'vehicle','person':'person'}  
    res['name'] = res['name'].map(rn)  
  
  
[]:  
    def midpoint(xmin, ymin, xmax, ymax):  
        center_w = xmax - xmin  
        center_h = ymax - ymin  
        center_x = 0.5*(xmin + xmax)  
        center_y = 0.5*(ymin + ymax)  
        return (center_x, center_y)  
    center = []  
    for i in range(4):  
        center.append(midpoint(res['xmin'][i], res['ymin'][i], res['xmax'][i], res['ymax'][i]))  
    distance = pow(pow((center[0][0]-center[1][0]),2) + pow((center[0][1]-center[1][1]),2),0.5)
```

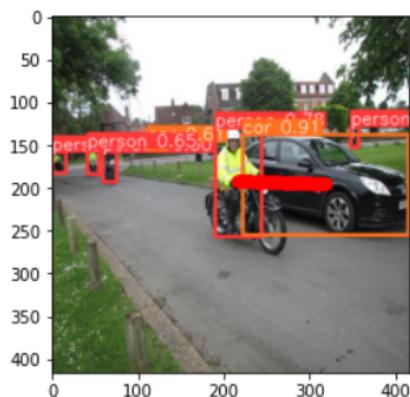
Plotting the difference

The distance between objects is plot in a form of red line

```
image = np.squeeze(results.render())
window_name = 'Image'
start_point = (int(center[0][0]), int(center[0][1]))

end_point = (int(center[1][0]), int(center[1][1]))

# Blue color in BGR
color = (255, 0, 0)
# Line thickness
thickness = 15
label = str(int(distance))
text_org = (int(center[0][0] + 200) , int(center[0][1]) - 50)
image = cv2.line(image, start_point, end_point, color, thickness)
image = cv2.putText(image, label, text_org, cv2.FONT_HERSHEY_SIMPLEX, 10, color, 12, cv2.LINE_AA)
# Displaying the image
plt.imshow(image)
plt.show()
```



Inferences of Study

- Annotations helps to data to know the coordinates to bounding boxes and object detection
- The pretrained model is used to handle data with high speed and better accuracy.

TASK 3

QUESTION 3

Download an image dataset of your choice for binary class classification. Perform the data augmentation techniques like flipping, rotation and transformation. Apply at least two object classification techniques both on the augmented as well as on the original dataset. Display the performance of the Algorithms. Prepare a comparison chart.

DATASET:

We have used cats and dogs dataset for this binary classification. Both train and testset consists of dogs and cats images.

Name	Date modified	Type	Size
test_set	11-12-2022 10:10	File folder	
training_set	11-12-2022 10:10	File folder	
augment_vgg18_cnn	11-12-2022 11:09	Jupyter Source File	779 KB

Models Used:

- Baseline Convolutional Neural Networks
 - Pretrained Model - VGG-16

BaseLine CNN:

Data Augmentation

CNN Model Fitting:

```
cnn = tf.keras.models.Sequential()  
cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation="relu", input_shape=[256,256,3])  
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2, padding='valid'))  
cnn.add(tf.keras.layers.Conv2D(filters=32, kernel_size=3, padding="same", activation="relu"))  
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2, padding='valid'))  
cnn.add(tf.keras.layers.Flatten())  
cnn.add(tf.keras.layers.Dense(units=128, activation='relu'))  
cnn.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))
```

1]

Python

```
cnn.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

2]

Python

```
cnn.summary()
```

3]

Python

VGG-16 Model Fitting:

```
base_model = VGG16(input_shape = IMG_SHAPE, include_top=False, weights='imagenet')
```

```
base_model.trainable = False
```

```
inputs = Input(shape = IMG_SHAPE)  
  
x = base_model(inputs, training = False)  
x = GlobalAveragePooling2D()(x)  
x = Dropout(0.5)(x)  
outputs = Dense(1, activation='sigmoid')(x)  
  
model = Model(inputs = inputs, outputs = outputs)
```

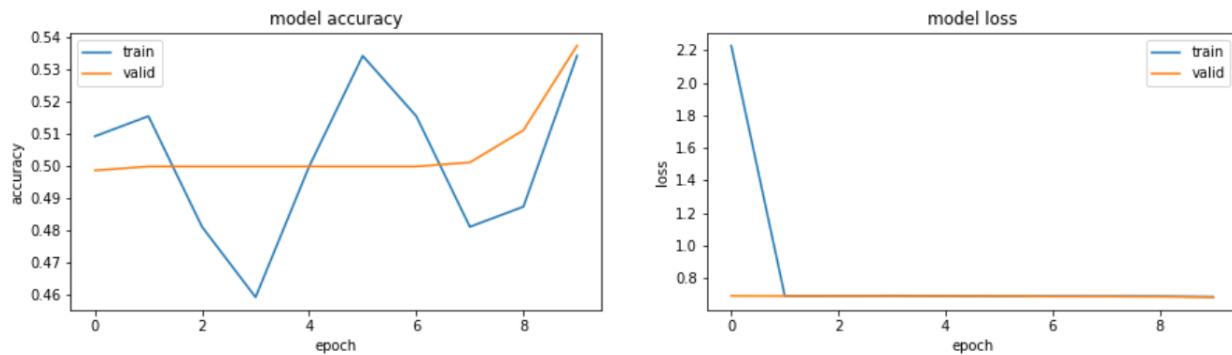
```
base_learning_rate = 0.001  
optimizer = Adam(learning_rate=base_learning_rate)
```

```
model.compile(  
    optimizer = optimizer,  
    loss = BinaryCrossentropy(),  
    metrics = ['accuracy'])
```

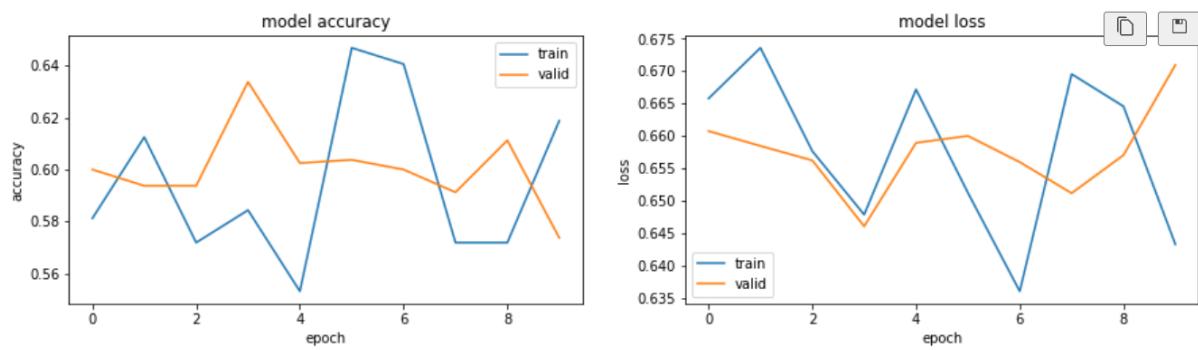
Results:

BaseLine CNN

Without Augmentation:

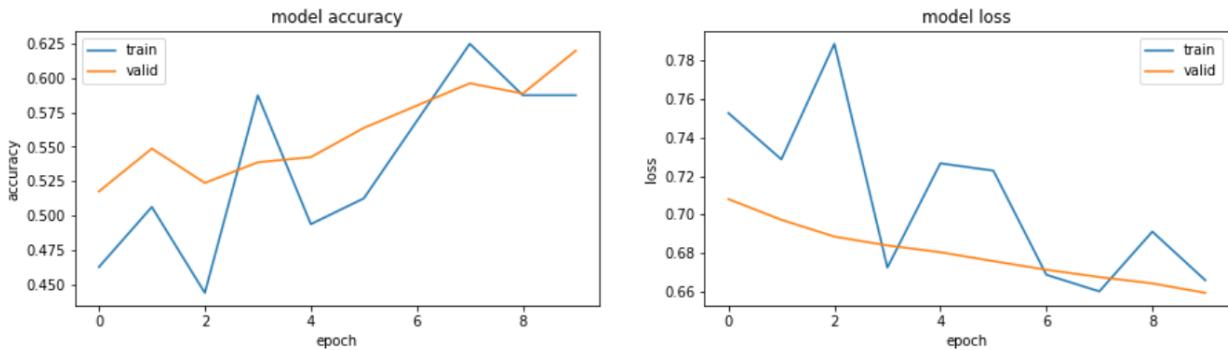


With Augmentation:

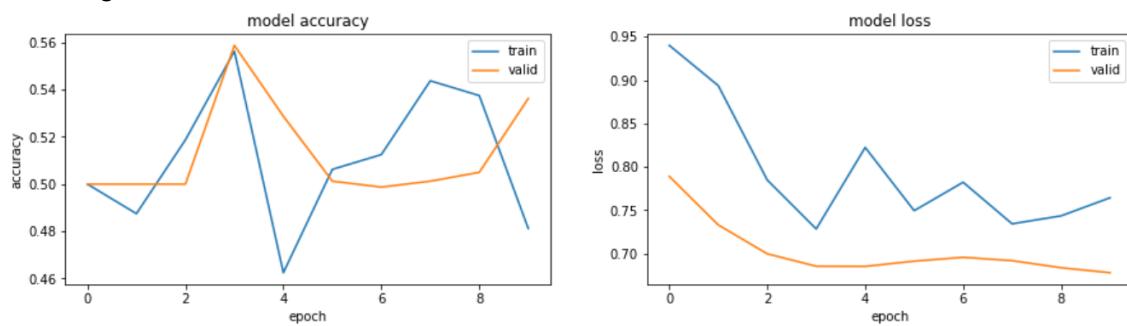


VGG-18

Without Augmentation:



With Augmentation:



Inferences of Study

- Augmentation Data for both baseline model and vgg-16 performs better than Non-Augmented Data.
- Pretrained Model along with augmented data performs better than baseline models.

TASK 4

QUESTION 4:

4. Collect images of vehicles with license plates written in Indian regional languages (eg. Hindi, Kannada, Tamil, Telugu, Bengali, etc.). Apply Image augmentation techniques on the collected images. Maintain separate folders for different language license plates. You may limit to 800 images in the dataset including the augmented images.

The aim is to get the license plates of different Indian languages and augment them and map them to the respective folders.

DATASET:

It consists of different license plates taken from different sites . I took 7 images for each class . Total of 35 images. It is saved in the JPEG format . It is stored separately in each folder.

Name		Date modified	Type
└── bengali		11-12-2022 16:07	File folder
└── hindi		09-12-2022 17:58	File folder
└── kannada		09-12-2022 17:42	File folder
└── tamil		11-12-2022 16:07	File folder
└── telugu		11-12-2022 16:07	File folder

Sdir gives the original path , augdir is the place to store the files. N - no of augmented images . I have given some additional features like img_size gives the (height and width) , color_mode is set to rgb by default. save _prefix is the prefix for the augmented images that will be saved for the new images .

Type of augmentation that is used on images

```
gen = ImageDataGenerator(rescale=1./255, rotation_range=15,
                        fill_mode='nearest', shear_range=0.2,
                        zoom_range=0.2, width_shift_range=0.2,
                        height_shift_range=0.2, horizontal_flip=True,
                        brightness_range=[0.5, 1.5])
```

Starting with an empty directory and creating a new one if the old path does not exist. If it already exists with some images it will delete them and create a new one , that prevents appending in the same files.

```

df=df.copy()
if os.path.isdir(augdir)
    shutil.rmtree(augdir)
os.mkdir(augdir)
for label in df['labels'].unique():
    classpath=os.path.join(augdir,label)
    os.mkdir(classpath)
total=0

```

Specifying no of images need to be created in the each folder and gives the total no of augmented images in each class.

```

In [5]: sdir=r'C:/Users/DELL/Desktop/main/license plate'
df=make_dataframe(sdir)
print (df.head())
print ('length of dataframe is ',len(df))

augdir=r'C:/Users/DELL/Desktop/main/new'
n=7
img_size=(224,224)
make_and_store_images(df, augdir, n, img_size, color_mode='rgb', save_prefix='aug-', save_format='jpg')

bengali      : 100%|██████████| 9/9 [00:00<00:00, 9022.16fil
es/s]
hindi       : 100%|██████████| 9/9 [00:00<00:00, 9418.35fil
es/s]
kannada     : 100%|██████████| 7/7 [00:00<00:00, 6983.86fil
es/s]
tamil        : 100%|██████████| 12/12 [00:00<?, ?fil
es/s]
telugu      : 100%|██████████| 7/7 [00:00<00:00, 17783.24fil
es/s]

filepaths  labels
0 C:/Users/DELL/Desktop/main/license plate\benga... bengali
1 C:/Users/DELL/Desktop/main/license plate\benga... bengali
2 C:/Users/DELL/Desktop/main/license plate\benga... bengali
3 C:/Users/DELL/Desktop/main/license plate\benga... bengali
4 C:/Users/DELL/Desktop/main/license plate\benga... bengali
length of dataframe is  44

```

Finally specify the path where to take the images and the path where to store and name of the file to be created . Here I have given new , it will create a folder name new and itself create the subfolders by taking the name form sdir path.

```

make_and_store_images(df, augdir, n, img_size, color_mode='rgb', save_prefix='aug-', save_format='jpg')

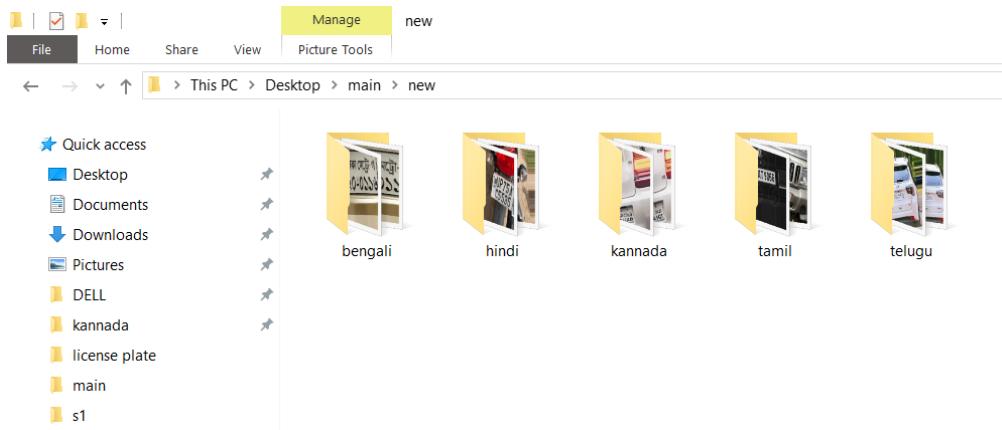
telugu      : 100%|██████████| 7/7 [00:00<?, ?fil
es/s]

filepaths  labels
0 C:/Users/DELL/Desktop/main/license plate\benga... bengali
1 C:/Users/DELL/Desktop/main/license plate\benga... bengali
2 C:/Users/DELL/Desktop/main/license plate\benga... bengali
3 C:/Users/DELL/Desktop/main/license plate\benga... bengali
4 C:/Users/DELL/Desktop/main/license plate\benga... bengali
length of dataframe is  35
Found 7 validated image filenames.      for class      bengali      creating 14  augmented images

Found 7 validated image filenames.      for class      hindi      creating 14  augmented images
Found 7 validated image filenames.      for class      kannada     creating 14  augmented images
Found 7 validated image filenames.      for class      tamil      creating 14  augmented images
Found 7 validated image filenames.      for class      telugu      creating 14  augmented images
Total Augmented images created= 70

```

OUTPUT:



BENGALI :



KANNADA:



HINDI:



TAMIL:



TELUGU:



TASK 5

QUESTION

Download the pytorch check point file from here (link to be added) and convert the file to .onnx

- Perform inferences on an onnx runtime session.
- Write a wrapper to perform the inference on video feed from webcam.
- Relevant papers to the .pth file:

Exporting a Model from PyTorch to ONNX and Running it using ONNX Runtime

Super-resolution model in PyTorch:

Super-resolution is a way of increasing the resolution of images, videos and is widely used in image processing or video editing.

```
▶ # Super Resolution model definition in PyTorch
import torch.nn as nn
import torch.nn.init as init

class SuperResolutionNet(nn.Module):
    def __init__(self, upscale_factor, inplace=False):
        super(SuperResolutionNet, self).__init__()

        self.relu = nn.ReLU(inplace=inplace)
        self.conv1 = nn.Conv2d(1, 64, (5, 5), (1, 1), (2, 2))
        self.conv2 = nn.Conv2d(64, 64, (3, 3), (1, 1), (1, 1))
        self.conv3 = nn.Conv2d(64, 32, (3, 3), (1, 1), (1, 1))
        self.conv4 = nn.Conv2d(32, upscale_factor ** 2, (3, 3), (1, 1), (1, 1))
        self.pixel_shuffle = nn.PixelShuffle(upscale_factor)

        self._initialize_weights()

    def forward(self, x):
        x = self.relu(self.conv1(x))
        x = self.relu(self.conv2(x))
        x = self.relu(self.conv3(x))
        x = self.pixel_shuffle(self.conv4(x))
        return x

    def _initialize_weights(self):
        init.orthogonal_(self.conv1.weight, init.calculate_gain('relu'))
        init.orthogonal_(self.conv2.weight, init.calculate_gain('relu'))
        init.orthogonal_(self.conv3.weight, init.calculate_gain('relu'))
        init.orthogonal_(self.conv4.weight)

# Create the super-resolution model by using the above model definition.
torch_model = SuperResolutionNet(upscale_factor=3)
```

Training Super-resolution Model using Pretrained-weight

```
▼ Load pretrained model weights

[4]
model_url = 'https://s3.amazonaws.com/pytorch/test_data/export/superres_epoch100-44c6958e.pth'
batch_size = 1 # just a random number

# Initialize model with the pretrained weights
map_location = lambda storage, loc: storage
if torch.cuda.is_available():
    map_location = None
torch_model.load_state_dict(model_zoo.load_url(model_url, map_location=map_location))

# set the model to inference mode
torch_model.eval()

Downloading: "https://s3.amazonaws.com/pytorch/test_data/export/superres_epoch100-44c6958e.pth"
100% [██████████] 234k/234k [00:00<0:00, 251kB/s]

SuperResolutionNet(
    (relu): ReLU()
    (conv1): Conv2d(1, 64, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))
    (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv3): Conv2d(64, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (conv4): Conv2d(32, 9, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (pixel_shuffle): PixelShuffle(upscale_factor=3)
)
```

Exporting a Model from PyTorch to ONNX:

```
[5] # Input to the model
x = torch.randn(batch_size, 1, 224, 224, requires_grad=True)
torch_out = torch_model(x)

torch.onnx.export(torch_model, # model being run
                 x, # model input (or a tuple for multiple inputs)
                 "super_resolution.onnx", # where to save the model (can be a file or file-like object)
                 export_params=True, # store the trained parameter weights inside the model file
                 opset_version=10, # the ONNX version to export the model to
                 do_constant_folding=True, # whether to execute constant folding for optimization
                 input_names = ['input'], # the model's input names
                 output_names = ['output'], # the model's output names
                 dynamic_axes={'input' : {0 : 'batch_size'}, # variable length axes
                               'output' : {0 : 'batch_size'}})
```

Now let's compute the output using ONNX Runtime's Python APIs. This part can normally be done in a separate process or on another machine, but we will continue in the same process so that we can verify that ONNX Runtime and PyTorch are computing the same value for the network.

In order to run the model with ONNX Runtime, we need to create an inference session for the model with the chosen configuration parameters (here we use the default config). Once the session is created, we evaluate the model using the run() api. The output of this call is a list containing the outputs of the model computed by ONNX Runtime.

Inferences on an onnx runtime session.

▼ Inferences on an onnx runtime session.

```
[11] import onnxruntime  
  
ort_session = onnxruntime.InferenceSession("super_resolution.onnx")  
  
def to_numpy(tensor):  
    return tensor.detach().cpu().numpy() if tensor.requires_grad else tensor.cpu().numpy()  
  
# compute ONNX Runtime output prediction  
ort_inputs = {ort_session.get_inputs()[0].name: to_numpy(x)}  
ort_outs = ort_session.run(None, ort_inputs)  
  
# compare ONNX Runtime and PyTorch results  
np.testing.assert_allclose(to_numpy(torch_out), ort_outs[0], rtol=1e-03, atol=1e-05)  
  
print("Exported model has been tested with ONNXRuntime, and the result looks good!")
```

Exported model has been tested with ONNXRuntime, and the result looks good!

Running the model on an image using ONNX Runtime

▼ Running the model on an image using ONNX Runtime

```
[13] from PIL import Image  
import torchvision.transforms as transforms  
  
img = Image.open("A-Cat.jpg")  
  
resize = transforms.Resize([224, 224])  
img = resize(img)  
  
img_ycbcr = img.convert('YCbCr')  
img_y, img_cb, img_cr = img_ycbcr.split()  
  
to_tensor = transforms.ToTensor()  
img_y = to_tensor(img_y)  
img_y.unsqueeze_(0)  
  
tensor([[[[0.7725, 0.7843, 0.7804, ..., 0.6824, 0.6784, 0.6902],  
         [0.7882, 0.7843, 0.7725, ..., 0.6549, 0.6667, 0.6824],  
         [0.8039, 0.7961, 0.7843, ..., 0.6588, 0.6706, 0.6784],  
         ...,  
         [0.7333, 0.7294, 0.7255, ..., 0.1020, 0.0941, 0.0902],  
         [0.7333, 0.7333, 0.7333, ..., 0.1020, 0.0902, 0.0863],  
         [0.7333, 0.7373, 0.7333, ..., 0.0980, 0.0863, 0.0824]]]])  
  
[19] ort_inputs = {ort_session.get_inputs()[0].name: to_numpy(img_y)}  
ort_outs = ort_session.run(None, ort_inputs)  
img_out_y = ort_outs[0]
```

Comparison between input and output data:

```
Comparision between input and output
```

```
[20] display(final_img) #final_img
```



```
[21] display(img) #input image
```

